

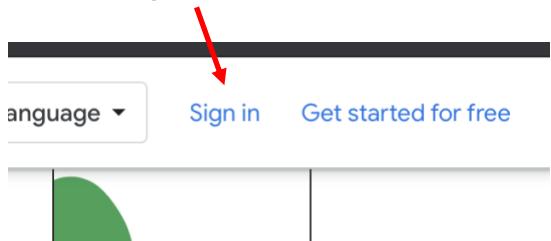
# Intro to DevOps, Containers & Orchestration



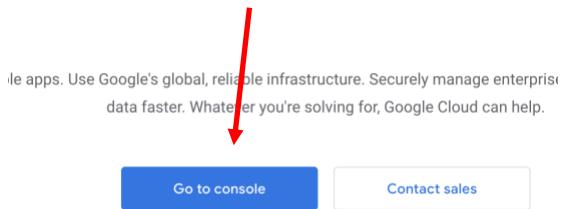
# Lab Section A – Access Google Cloud Platform

Login to GCP - <https://cloud.google.com>

Click “Sign In” button

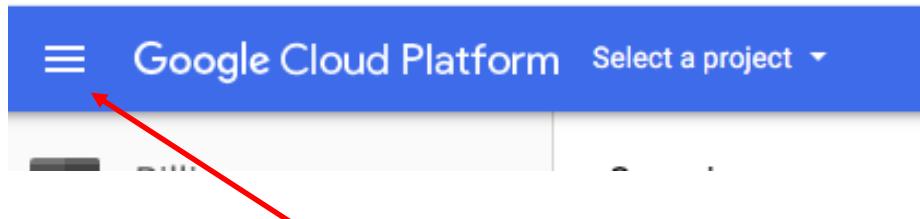


Click Go to console



Select an existing project or create a new one:

A screenshot of the Google Cloud Platform dashboard. On the left, there is a sidebar with "Project info" details: Project name (My First Project), Project ID (marine-order-267121), and Project number (901212451927). In the center, a modal window titled "Select a project" is open, showing a search bar and a list of recent projects. The "RECENT" tab is selected, and "My First Project" is listed with a checkmark next to it. A red arrow points from the text above to the "NEW PROJECT" button in the top right of the modal. Another red arrow points to the "My First Project" entry in the list.

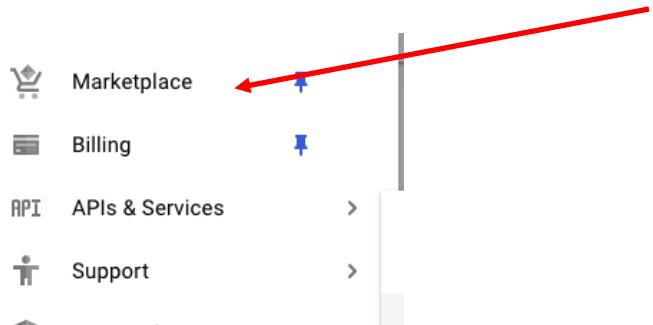


Remember this ‘hamburger’ icon – it brings up the main menu in GCP

**IMPORTANT NOTE:** If you quit lab early be sure to delete all compute instances, Kubernetes Clusters and Network Load Balancers – the clock is running...

## Lab Section B – Installing Docker

Click the 3 lines (hamburger) on menu and choose “Marketplace”



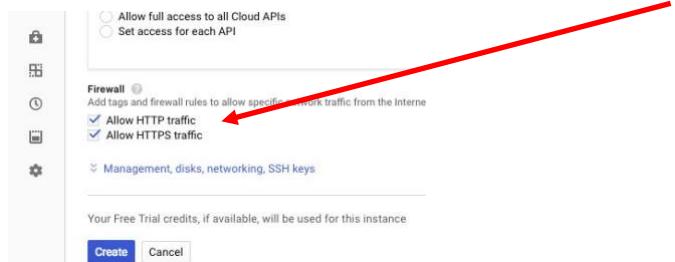
type “Centos” in the search bar and double click “Centos 7”



Click “Launch”

LAUNCH

Accept all defaults (or you can choose region local to you if you like)  
**AND** click Allow http and https traffic under firewall section



Before clicking “Create” View REST and/or command line to view REST or CLI output  
(if you had preferred using one of these options)

Create

Cancel

Equivalent REST or command line

Now click create

Once started (icon will turn green) click on the SSH tab to open a remote console

VM Instances					
<input type="checkbox"/>	Name ^	Zone	Recommendation	Internal IP	External IP
<input type="checkbox"/>	<input checked="" type="checkbox"/> centos-7-1	us-east1-b		10.142.0.2	104.196.105.65 



**NOTE:** typed commands are shown in *italics* and highlighted. Some commands are in a smaller font (on one single line) as a work-around with PDF line break issues

### Install the docker engine in the Centos VM

*sudo yum check-update*  
Update the package database

*curl -fsSL <https://get.docker.com> / sh*  
- Add official Docker repository, download and install latest Docker version

*sudo systemctl start docker*  
- Start the Docker daemon

*sudo systemctl status docker*  
- Verify Docker is running

You should see message similar to the following –

```
[klowenst@centos-7-1 ~]$ sudo systemctl status docker
● docker.service - Docker Application Container Engine
  Loaded: loaded (/usr/lib/systemd/system/docker.service; disabled; vendor
  Active: active (running) since Thu 2020-10-29 13:41:32 UTC; 11s ago
    Docs: https://docs.docker.com
  Main PID: 1636 (dockerd)
     Tasks: 10
```

*sudo systemctl enable docker.service*  
- Start the Docker daemon on start-up

*sudo usermod -a -G docker \$USER*  
- Run Docker daemon as “non-sudo” user

Exit the console by typing “exit”

Reopen the console to enable docker commands to run as a non-sudo user

type

*docker info*

- Verify docker daemon is running as regular user

Running command as non-sudo user should return status similar to below

```
[klowenst@centos-7-1 ~]$ docker info
Client:
  Debug Mode: false

Server:
  Containers: 0
    Running: 0
    Paused: 0
    Stopped: 0
  Images: 1
  Server Version: 19.03.13
  Storage Driver: overlay2
    Backing Filesystem: xfs
    Supports d_type: true
    Native Overlay Diff: true
  Logging Driver: json-file
  Cgroup Driver: cgroupfs
  Plugins:
    Volume: local
    Network: bridge host ipvlan macvlan null overlay
    Log: awslogs fluentd gcplogs gelf journald json-file
```

Extra credit: If you would like to have the ability to see docker command options after typing ‘docker’ then tab key – follow these instructions.

```
sudo yum install -y bash-completion bash-completion-extras
```

```
sudo curl -L https://raw.githubusercontent.com/docker/compose/1.22.0/contrib/completion/bash/docker-compose -o /etc/bash_completion.d/docker-compose
```

```
source ~/.bash_profile
```

```
exec $SHELL -l
```

## Working with Images

From the SSH window enter the following commands

*docker images*

- Note that there are currently no images on the system

*docker ps*

- There are no running containers

*docker pull ubuntu*

- download an image

```
[klowenst@centos-7-docker-installed ~]$ docker pull ubuntu
Using default tag: latest
latest: Pulling from library/ubuntu
6a5697faee43: Pull complete
ba13d3bc422b: Pull complete
a254829d9e55: Pull complete
Digest: sha256:fff16ee1a8ae92867721d90c59a75652ea66d29c05294e6e2f898704bdb8cf1
Status: Downloaded newer image for ubuntu:latest
docker.io/library/ubuntu:latest
[klowenst@centos-7-docker-installed ~]$ █
```

Docker is pulling the layered image

type *docker images* command again

```
[klowenst@centos-7-1 ~]$ docker images
REPOSITORY      TAG      IMAGE ID      CREATED      SIZE
ubuntu          latest   d70eaf7277ea  5 days ago   72.9MB
[klowenst@centos-7-1 ~]$ █
```

We have the “latest” Ubuntu image stored locally, - we could have specified a version.

type *docker inspect <image ID>* - copy and paste <image ID> from previous “docker images” output

- i.e. (from above screen shot) – docker inspect d70eaf7277ea

## Examine the output

```
[klowenst@centos-7-docker-installed ~]$ cat /etc/docker/daemon.json
{
    "registry-mirrors": [
        "https://mirrors.aliyun.com/dockerhub/"
    ],
    "log-driver": "json-file",
    "log-opts": {
        "max-size": "100m"
    },
    "storage-driver": "overlay2",
    "storage-opts": {
        "data-root": "/var/lib/docker"
    },
    "dns": [
        "8.8.8.8",
        "8.8.4.4"
    ],
    "dns-search": [
        "centos"
    ],
    "ip-masq": true,
    "ulimits": {
        "nofile": 65536
    },
    "oom-scoreadj": -1000,
    "cpu-shares": 1024,
    "cpu-period": 1000000,
    "cpu-quota": 1000000,
    "cpuacct-group": "cpuacct"
}
[klowenst@centos-7-docker-installed ~]$
```

Type

```
ps -aux
```

Notice all the processes currently running on your Centos 7 VM OS

```
root      1006  0.0  0.2 218548  7540 ?      Ssl  14:30  0:00 /usr/sbin/rsyslogd -n
root     1248  0.0  0.3 214788 12700 ?      Ss  14:30  0:00 /usr/bin/python /usr/bin/google_network_daemon
root     1250  0.0  0.3 217320 13156 ?      Ss  14:30  0:00 /usr/bin/python /usr/bin/google_accounts_daemon
root     1254  0.0  0.3 214776 12624 ?      Ss  14:30  0:00 /usr/bin/python /usr/bin/google_clock_skew_daemon
root     1256  0.0  0.1 112920  4344 ?      Ss  14:30  0:00 /usr/sbin/sshd -D
root     1313  0.0  0.0  89700  2084 ?      Ss  14:30  0:00 /usr/libexec/postfix/master -w
postfix   1316  0.0  0.1  89804  4056 ?      S  14:30  0:00 pickup -l -t unix -u
postfix   1317  0.0  0.1  89872  4080 ?      S  14:30  0:00 qmgr -l -t unix -u
root     1716  0.0  1.0 489976 37764 ?      Ssl 14:36  0:00 /usr/bin/containerd
root     1717  0.1  2.4 570244  88212 ?      Ssl 14:36  0:02 /usr/bin/dockerd -H fd:// --containerd=/run/contai
root     1904  0.0  0.1 158920  5756 ?      Ss  14:37  0:00 sshd: kubefy2020 [priv]
kubefy2+  1907  0.0  0.0 158920  2436 ?      S  14:37  0:00 sshd: kubefy2020@pts/0
kubefy2+  1908  0.0  0.1 117072  3708 pts/0      Ss  14:37  0:00 /bin/bash -l
kubefy2+  1909  0.0  0.0  72248  2676 ?      Ss  14:37  0:00 /usr/libexec/openssh/sftp-server
root     2054  0.0  0.0      0  0 ?      S  14:45  0:00 [kworker/0:2]
root     2139  0.0  0.0      0  0 ?      R  14:50  0:00 [kworker/0:1]
root     2362  0.0  0.0      0  0 ?      S  14:55  0:00 [kworker/0:0]
root     2448  0.0  0.1 112920  4284 ?      Ss  14:58  0:00 sshd: [accepted]
kubefy2+  2449  0.0  0.0 155372  1864 pts/0      R+ 14:58  0:00 ps -aux
[kubefy2020@centos-7-1 ~]$
```

Now type

```
cat /etc/os-release
```

```
[kubefy2020@centos-7-1-docker ~]$ cat /etc/os-release
NAME="CentOS Linux"
VERSION="7 (Core)"
ID="centos"
ID_LIKE="rhel fedora"
VERSION_ID="7"
PRETTY_NAME="CentOS Linux 7 (Core)"
```

Notice you are running “CentOS Linux 7”

## Creating the first container

type

```
docker run -it ubuntu
```

(-I = interactive –t=terminal)

you are now “in” the container – notice the prompt change

Type

```
ps -aux
```

```
[klowenst@centos-7-1 ~]$ docker run -it ubuntu
root@67ace5d20aec:/# ps -aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root         1  0.1  0.0    4100  2164 pts/0    Ss   13:49   0:00 /bin/bash
root         8  0.0  0.0    5880  1408 pts/0    R+   13:49   0:00 ps -aux
```

Only 2 processes are running in this container – “/bin/bash” & the command you just entered, “ps aux” (notice you are running as root user inside container)

Now type

```
cat /etc/os-release again
```

```
root@67ace5d20aec:/# cat /etc/os-release
NAME="Ubuntu"
VERSION="20.04.1 LTS (Focal Fossa)"
ID=ubuntu
ID_LIKE=debian
PRETTY_NAME="Ubuntu 20.04.1 LTS"
VERSION_ID="20.04"
HOME_URL="https://www.ubuntu.com/"
SUPPORT_URL="https://help.ubuntu.com/"
```

We are now running “Ubuntu 20.04” inside the container

Let’s exit the container but leave it running

Press and hold “CTRL” key then “p-q”

You should now be back at your linux prompt

type:

*docker ps*

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
NAMES				
67ace5d20aec	ubuntu	"/bin/bash"	4 minutes ago	Up 4 minutes
[klowenst@centos-7-1 ~]\$				

Notice that the container is running, has a unique ID, the command is the “/bin/bash/” shell and a random short name has been assigned

Reattach to the running container

*docker attach <image ID>*

- where <image ID> is container ID from previous docker ps command
- we are now back in the container (notice prompt)

While back in the container type

*traceroute yahoo.com*

the command will fail as traceroute package is not installed

Let's add traceroute to this container

*apt-get update*

*apt-get install traceroute*

now let's try the command again

*traceroute yahoo.com*

Exit the container by typing “exit”

type

*docker ps*

notice the container is no longer running

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
NAMES				
[klowenst@centos-7-1 ~]\$				
[klowenst@centos-7-1 ~]\$				

Now we can create a new image, adding the new “layer” (with traceroute installed)

Let's get the container ID (the appended container still exists but not running)

`docker ps -a` (-a' signifies all, not just running containers)

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
67ace5d20aec	ubuntu	"/bin/bash"	7 minutes ago	Exited

copy the container ID and paste after commit

`docker commit <container ID> ubuntu_traceroute`

create a new image

type `docker images` command and we will now see 2 images

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
ubuntu_traceroute	latest	2c5fca0d2e83	20 seconds ago	99.7MB
ubuntu	latest	d70eaf7277ea	5 days ago	72.9MB

Note: You can restart either container by using a “`docker run -it <container ID OR image name>`“

Type `docker inspect ubuntu`

We see the original 3 layers

```
},  
"RootFS": {  
    "Type": "layers",  
    "Layers": [  
        "sha256:47dde53750b4a8ed24acebe52cf31ad131e73a9611048fc2f92c9b6274ab4bf3",  
        "sha256:0c2689e3f9206b1c4adfb16a1976d25bd270755e734588409b31ef29e3e756d6",  
        "sha256:cc9d18e90faad04bc3893cf5aa50b7846ee75f48f5b8377a213fa52af2189095c"  
    ]  
}
```

Type `docker inspect ubuntu_traceroute`

We now see the original 3 layers plus the layer we added

```
"RootFS": {  
    "Type": "layers",  
    "Layers": [  
        "sha256:47dde53750b4a8ed24acebe52cf31ad131e73a9611048fc2f92c9b6274ab4bf3",  
        "sha256:0c2689e3f9206b1c4adfb16a1976d25bd270755e734588409b31ef29e3e756d6",  
        "sha256:cc9d18e90faad04bc3893cf5aa50b7846ee75f48f5b8377a213fa52af2189095c",  
        "sha256:f5d8b04f08578a87284c7e04c3958c25702bfff99ed2218e7d676b28d291f0421"  
    ]  
}
```

## Running a Web Server

type

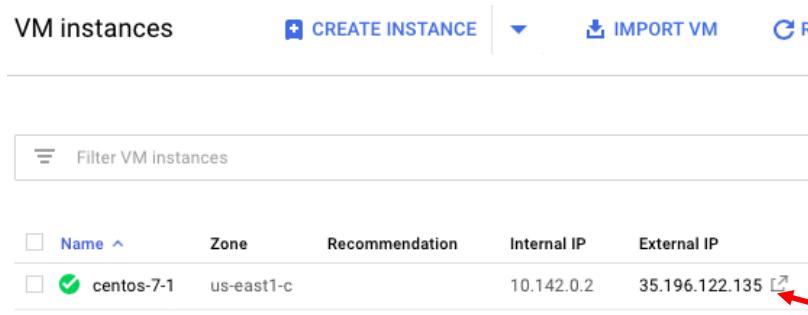
```
docker run -d -p 80:80 nginx
```

- docker will pull and run the image (-d = in detached mode / -p = publish container's port(s) to the host).
- You can verify by typing "docker ps"

Go to the Compute Engine VM Instances page

- (Click on 3 bars on top of page  > then click "Compute Engine")

Get the external IP address of the server.



<input type="checkbox"/>	Name	Zone	Recommendation	Internal IP	External IP	
<input checked="" type="checkbox"/>	centos-7-1	us-east1-c		10.142.0.2	35.196.122.135	

In a new browser type in `http://<external IP address>`

The NGINX default web page should show.



## Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to [nginx.org](http://nginx.org)  
Commercial support is available at [nginx.com](http://nginx.com)

Let's stop this nginx container and delete it.

Back in the Centos VM – type

```
docker stop <container id>
```

Now enter `docker ps` to verify it is no longer running

## Creating Dockerfiles

Will a create simple Dockerfile that uses nginx and copies content into the container

First let's install a text editor (so we don't have to struggle using 'vi')

```
sudo yum install nano -y
```

Let's get some content

type

```
curl http://www.gutenberg.org/files/2600/2600-h/2600-h.htm >index.html
```

(gutenberg.org may have occasional outages as improvements are made. If you run into a 'time-out' try one of these links instead)

```
curl http://dev.gutenberg.org/files/2600/2600-h/2600-h.htm >index.html
```

```
curl http://www.authorama.com/alice-in-wonderland-1.html >index.html
```

type

```
cat index.html
```

- to verify contents are there (this is a large file, hit CTRL-C to stop)

Let's create a docker file

type

```
nano Dockerfile ('Dockerfile' is case sensitive)
```

Paste the following 2 lines into the file

```
from nginx
copy ./index.html /usr/share/nginx/html
```

```
from nginx
copy ./index.html /usr/share/nginx/html
█
```

To save file

hit ‘control – o’ and ‘enter” to write file

hit ‘control – x’ to exit nano

type

```
cat Dockerfile
```

- to verify file was written

we will now build the image

type

```
docker build -t testweb .
```

- (be sure to include “.” ( . = this directory)

Confirm that the image was built

type

```
docker images
```

- to verify image (testweb) is there

```
[klowenst@centos-7-1 ~]$ docker images
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
testweb              latest   ecced2a3e8b3  6 seconds ago  137MB
ubuntu_traceroute   latest   2c5fca0d2e83  3 minutes ago  99.7MB
ubuntu               latest   d70eaf7277ea  5 days ago    72.9MB
nginx                latest   f35646e83998  2 weeks ago   133MB
[klowenst@centos-7-1 ~]$ █
```

type

*docker history testweb*

- to see the layer we added to image several seconds ago

IMAGE	CREATED	CREATED BY	SIZE
<b>COMMENT</b>			
ecced2a3e8b3	45 seconds ago	/bin/sh -c #(nop) COPY file:ea63ff61269ed737...	4.07MB
f35646e83998	2 weeks ago	/bin/sh -c #(nop) CMD ["nginx" "-g" "daemon...]	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) STOPSIGNAL SIGTERM	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) EXPOSE 80	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) ENTRYPOINT ["/docker-entr...]	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) COPY file:0fd5fca330dcd6a7...	1.04kB
<missing>	2 weeks ago	/bin/sh -c #(nop) COPY file:1d0a4127e78a26c1...	1.96kB
<missing>	2 weeks ago	/bin/sh -c #(nop) COPY file:e7e183879c35719c...	1.2kB
<missing>	2 weeks ago	/bin/sh -c #(nop) COPY file:63.6MB	63.6MB
<missing>	2 weeks ago	/bin/sh -c set -x && addgroup --system ...	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) ENV PKG_RELEASE=1~buster	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) ENV NJS_VERSION=0.4.4	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) ENV NGINX_VERSION=1.19.3	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) LABEL maintainer=NGINX Do...	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) CMD ["bash"]	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) ADD file:0dc53e7886c35bc21...	69.2MB

Type *docker history nginx* to view original image we copied from

- notice there is one less “layer”

IMAGE	CREATED	CREATED BY	SIZE
<b>COMMENT</b>			
f35646e83998	2 weeks ago	/bin/sh -c #(nop) CMD ["nginx" "-g" "daemon...]	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) STOPSIGNAL SIGTERM	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) EXPOSE 80	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) ENTRYPOINT ["/docker-entr...]	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) COPY file:0fd5fca330dcd6a7...	1.04kB
<missing>	2 weeks ago	/bin/sh -c #(nop) COPY file:1d0a4127e78a26c1...	1.96kB
<missing>	2 weeks ago	/bin/sh -c #(nop) COPY file:e7e183879c35719c...	1.2kB
<missing>	2 weeks ago	/bin/sh -c #(nop) set -x && addgroup --system ...	63.6MB
<missing>	2 weeks ago	/bin/sh -c #(nop) ENV PKG_RELEASE=1~buster	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) ENV NJS_VERSION=0.4.4	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) ENV NGINX_VERSION=1.19.3	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) LABEL maintainer=NGINX Do...	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) CMD ["bash"]	0B
<missing>	2 weeks ago	/bin/sh -c #(nop) ADD file:0dc53e7886c35bc21...	69.2MB

Now run the container

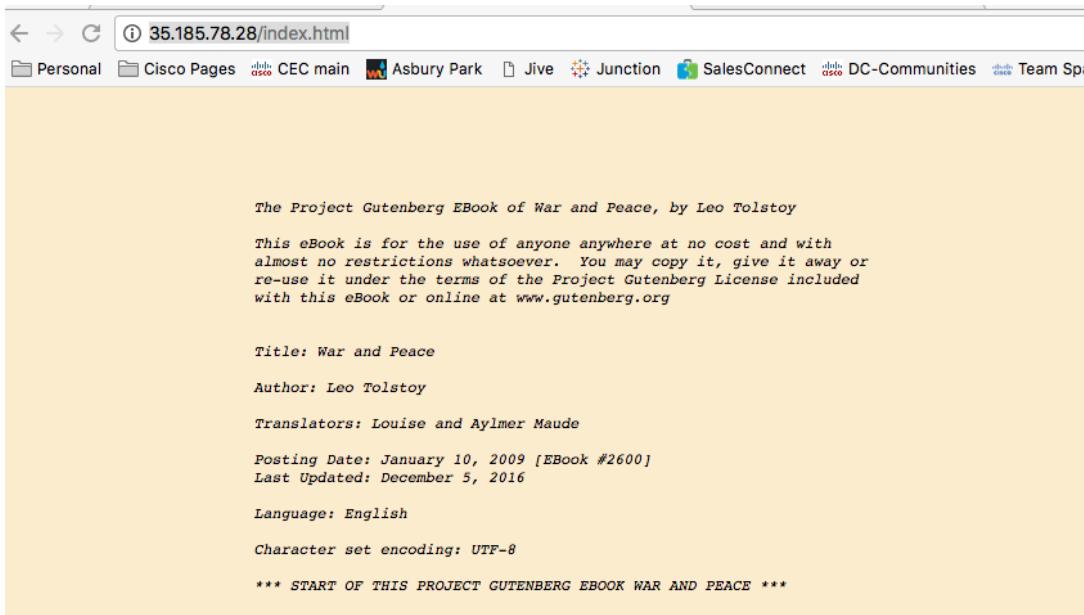
*docker run -d -p 80:80 testweb*

Retrieve the IP address of VM from GCP console (VM instances)

A screenshot of the Google Cloud Platform Compute Engine VM instances page. On the left, there's a sidebar with options: Compute Engine (selected), VM instances (highlighted in blue), Instance groups, Instance templates, Disks, and Snapshots. The main area shows a table of VM instances. A red arrow points to the 'External IP' column for the 'centos-7-1' row, which contains the value '35.185.78.28'.

Verify web server is running by opening browser with this address - adding  
"/index.html"      [http://<your\\_external\\_ip>/index.html](http://35.185.78.28/index.html)

**NOTE: If you have issues viewing link – try another browser or clear the cache and try again.**



```
type
docker ps
docker stop <container_ID>
```

# STOP HERE

# Lab Section C – Running Kubernetes

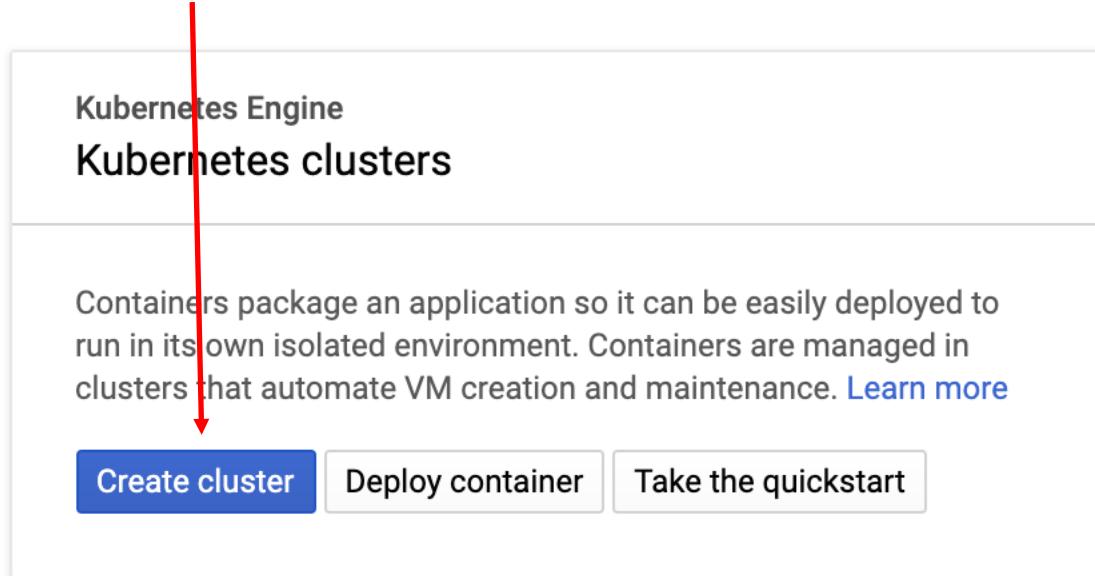
## Building your Cluster

Within your project

> click on the 3 horizontal bars to the left of Google Could Platform 

> click on Kubernetes Engine (wait for it to initialize – you may have to refresh screen)

> click “Create cluster”



Change to **your** zone

Location type  
 Zonal  
 Regional

Zone  ▼ ?

Specify default node locations ?

Current default: us-east1-b

### Master version

Choose a release channel for automatic management of your cluster's version and upgrade cadence. Choose a static version for more direct management of your cluster's version. [Learn more](#).

Static version

Release channel

Static version  ▼

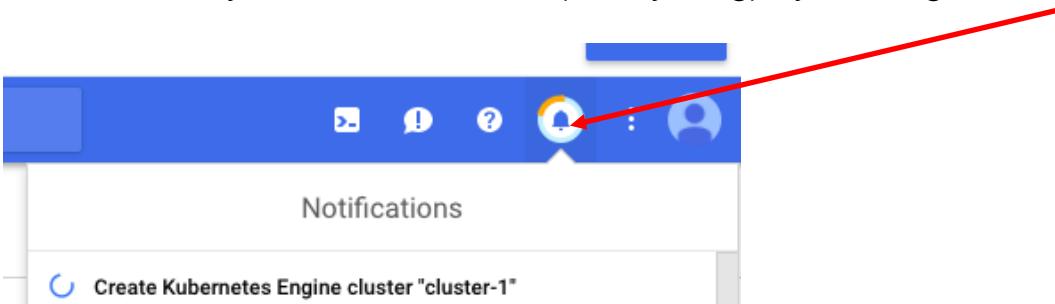
Choose **1.16.13-gke.401 (default)** cluster version  
(note: this version may vary as default is updated often)

View REST and/or command line to view REST or CLI output



Click "Create"

You can always check the status (of anything) by clicking the Notifications icon



When cluster icon turns green - click connect

Name	Location	Cluster size	Total cores	Total memory	Notifications	Labels	Connect
cluster-1	us-central1-a	3	3 vCPUs	11.25 GB			

A dialog box will open –

Click "Run in Cloud Shell" – The cloud shell console will open and populate the command – hit "enter" to continue

Connect to the cluster

You can connect to your cluster via command line or using a dashboard.

#### Command-line access

Configure `kubectl` command line access by running the following command:

```
$ gcloud container clusters get-credentials standard-cluster-1 --zone us-central1-a -
```

Run in Cloud Shell

Run the following commands to verify cluster is up and running

*kubectl cluster-info*

```
0 upgraded, 0 newly installed, 0 to remove and 17 not upgraded.  
kubefy2020@cloudshell:~ (marine-order-267121)$ kubectl cluster-info  
Kubernetes master is running at https://35.231.41.175  
GLBCDefaultBackend is running at https://35.231.41.175/api/v1/namespaces/kube-system/services/default-http-backend:http/proxy  
Heapster is running at https://35.231.41.175/api/v1/namespaces/kube-system/services/heapster/proxy  
KubeDNS is running at https://35.231.41.175/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy  
Metrics-server is running at https://35.231.41.175/api/v1/namespaces/kube-system/services/https:metrics-server:/proxy
```

*kubectl get nodes*

```
kubefy2020@cloudshell:~ (marine-order-267121)$ kubectl get nodes  
NAME STATUS ROLES AGE VERSION  
gke-standard-cluster-1-default-pool-9cb2ae4a-1j6z Ready <none> 23m v1.13.11-gke.23  
gke-standard-cluster-1-default-pool-9cb2ae4a-wsr6 Ready <none> 23m v1.13.11-gke.23  
gke-standard-cluster-1-default-pool-9cb2ae4a-xngg Ready <none> 23m v1.13.11-gke.23  
kubefy2020@cloudshell:~ (marine-order-267121)$ █
```

*kubectl get pods --all-namespaces*

NAME	READY	STATUS	ROLES	AGE	VERSION
kube-system/fluentd-gcp-v3.2.0-jpvrd	2/2	Running	0	24m	
kube-system/fluentd-gcp-v3.2.0-m85hz	3/3	Running	0	24m	
kube-system/heapster-7f47454698-dtgct	4/4	Running	0	25m	
kube-system/kube-dns-79868f54c5-49nzm	4/4	Running	0	24m	
kube-system/kube-dns-79868f54c5-zrpvq	1/1	Running	0	24m	
kube-system/kube-dns-autoscaler-bb58c6784-xt2tj	1/1	Running	0	24m	
kube-system/kube-proxy-gke-standard-cluster-1-default-pool-9cb2ae4a-1j6z	1/1	Running	0	24m	
kube-system/kube-proxy-gke-standard-cluster-1-default-pool-9cb2ae4a-wsr6	1/1	Running	0	24m	
kube-system/kube-proxy-gke-standard-cluster-1-default-pool-9cb2ae4a-xngg	1/1	Running	0	24m	
kube-system/l7-default-backend-fd59995cd-cjzms	1/1	Running	0	25m	
kube-system/metrics-server-v0.3.1-57c75779f-ljx2n	2/2	Running	0	24m	
kube-system/prometheus-to-sd-67fdw	2/2	Running	0	24m	
kube-system/prometheus-to-sd-6j4gj	2/2	Running	0	24m	
kube-system/prometheus-to-sd-vcvqg	2/2	Running	0	24m	

```
kubefy2020@cloudshell:~ (marine-order-267121)$ █
```

## Exposing an External IP Address to Access an Application in a Cluster

`kubectl apply -f https://k8s.io/examples/service/load-balancer-example.yaml`

- Create a Hello World application in your cluster

`load-balancer-example.yaml` (from hyperlink above)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app.kubernetes.io/name: load-balancer-example
    name: hello-world
spec:
  replicas: 5
  selector:
    matchLabels:
      app.kubernetes.io/name: load-balancer-example
  template:
    metadata:
      labels:
        app.kubernetes.io/name: load-balancer-example
    spec:
      containers:
        - image: gcr.io/google-samples/node-hello:1.0
          name: hello-world
          ports:
            - containerPort: 8080
```

`kubectl get deployments hello-world`

`kubectl describe deployments hello-world`

- Display information about the Deployment

`kubectl get replicsets`

`kubectl describe replicsets`

- Display information about your ReplicaSet objects

`kubectl expose deployment hello-world --type=LoadBalancer --name=my-service`

- Create a Service object that exposes the deployment

`kubectl get services my-service`

- Display information about the Service

Output: (may take a minute or so to change from <pending> to external IP address)

You also have the option to append “watch” to above command. This will show you status of command every 2 seconds (CTRL-C to exit)

`watch kubectl get services my-service`

```
kubefy2020@cloudshell:~ (marine-order-267121)$
kubefy2020@cloudshell:~ (marine-order-267121)$ kubectl get services my-service
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
my-service   LoadBalancer   10.12.2.215   35.202.225.104   8080:31935/TCP   60s
kubefy2020@cloudshell:~ (marine-order-267121)$
```

*kubectl describe services my-service*

- Display detailed information about the Service

Output:

```
my-service   LoadBalancer   10.12.2.215   35.202.225.104   8080:31935/TCP   60s
kubefy2020@cloudshell:~ (marine-order-267121)$ kubectl describe services my-service
Name:           my-service
Namespace:      default
Labels:         app.kubernetes.io/name=load-balancer-example
Annotations:    <none>
Selector:       app.kubernetes.io/name=load-balancer-example
Type:          LoadBalancer
IP:            10.12.2.215
LoadBalancer Ingress: 35.202.225.104
Port:          <unset>  8080/TCP
TargetPort:    8080/TCP
NodePort:      <unset>  31935/TCP
Endpoints:     10.8.0.6:8080,10.8.1.4:8080,10.8.1.5:8080 + 2 more...
Session Affinity: None
```

Make note of the external IP address exposed by your service (LoadBalancer Ingress).  
Also note the port value (8080)

You can see that the service has several endpoints:

10.8.0.6:8080,10.8.1.4:8080,10.8.1.5:8080 + 2 more. These are internal addresses of the pods that are running the Hello World application. To verify these are pod addresses, enter this command:

*kubectl get pods --output=wide*

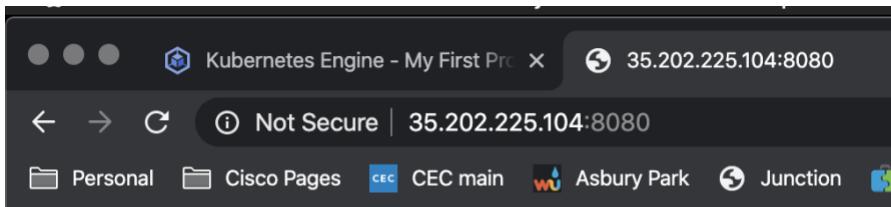
Output:

```
kubefy2020@cloudshell:~ (marine-order-267121)$
kubefy2020@cloudshell:~ (marine-order-267121)$ kubectl get pods --output=wide
NAME           READY   STATUS    RESTARTS   AGE     IP           NODE
hello-world-5885cc9fd8-54j6d  1/1    Running   0          6m9s   10.8.1.4   gke-standard-cluster-1-default-pool-ffa4414f-xj0s
hello-world-5885cc9fd8-97m4c  1/1    Running   0          6m9s   10.8.1.5   gke-standard-cluster-1-default-pool-ffa4414f-xj0s
hello-world-5885cc9fd8-cxn85  1/1    Running   0          6m9s   10.8.0.6   gke-standard-cluster-1-default-pool-ffa4414f-53fr
hello-world-5885cc9fd8-m8tr9  1/1    Running   0          6m9s   10.8.2.6   gke-standard-cluster-1-default-pool-ffa4414f-qnfz
hello-world-5885cc9fd8-w8g9n  1/1    Running   0          6m9s   10.8.2.7   gke-standard-cluster-1-default-pool-ffa4414f-qnfz
kubefy2020@cloudshell:~ (marine-order-267121)$
```

Open a browser and connect to: <http://<external-ip>:<port>>

- from kubectl get services my-service or kubectl describe services my-service

output should return “Hello Kubernetes !”



Hello Kubernetes!

## Create a Guestbook Kubernetes Engine cluster

We will demonstrate how to build a simple multi-tier web application (Guestbook with Redis and PHP) using Kubernetes Engine. The tutorial application is a guestbook that allows visitors to enter text in a log and to see the last few logged entries.

go to [tiny.cc/ourlab](http://tiny.cc/ourlab)

Download yaml.zip

Name	Size
yaml	

**yaml.zip**

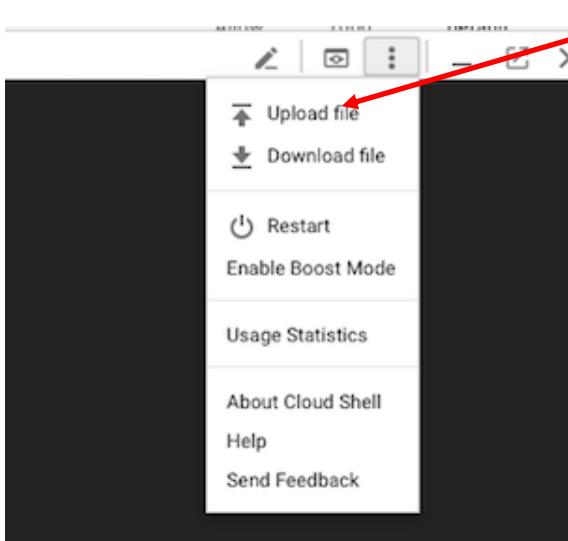
Details

Size  
6.66 KB

Sign in Sign up [Download](#)

Unzip the folder to your desktop

Upload yaml files to cloud shell (via cloud shell tool)



In Cloud Shell type

`ls -l`

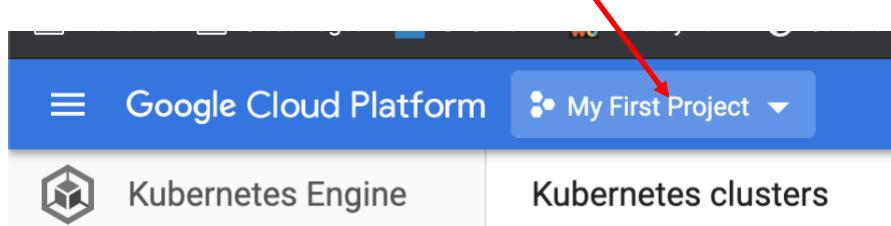
- in console to verify yaml files - then continue with lab

In your Cloud Shell console type:

`gcloud config set project xxxxxxxx`

- where “xxxxxxxx” is your project ID.

from GCP console click on Project Name



to view project ID

Name	ID
My First Project	marine-order-267121

A screenshot of the GCP Kubernetes Engine page. It shows a table with two columns: "Name" and "ID". There is one row in the table. A red arrow points from the text "to view project ID" down to the "ID" column of the table.

`gcloud config set compute/zone xxxxxxxx`

- where xxxxxxxx is your location (locate in GCP – Main Menu > Kubernetes Engine > location)

Filter by label or name			
<input type="checkbox"/>	Name	Location	Cluster size
<input checked="" type="checkbox"/>	standard-cluster-1	us-east1-b	3

A screenshot of the GCP Kubernetes Engine page showing a list of clusters. The first cluster, "standard-cluster-1", has a red circle drawn around its "Location" column value "us-east1-b".

## Create the Kubernetes Engine via CLI

```
gcloud container clusters create guestbook --num-nodes=3
```

(disregard any warnings)

Note: You should now notice 2 clusters in the GCP dashboard

<input type="checkbox"/>	Name ^	Location	Cluster size	Total cores	Total memory
<input type="checkbox"/>	guestbook	us-east1-b	3	3 vCPUs	11.25 GB
<input type="checkbox"/>	standard-cluster-1	us-east1-b	3	3 vCPUs	11.25 GB

Get a list of clusters in your project or details for a single cluster (via CLI)

```
gcloud container clusters list
```

- should see both clusters

```
gcloud container clusters describe guestbook
```

- get information on the guestbook cluster

### Step 1 – Set up a Redis Master

The guestbook application uses Redis to store its data. It writes its data to a Redis master instance and reads data from multiple Redis worker (slave) instances. The first step is to deploy a Redis master.

We will use the manifest yaml file named redis-master-deployment to deploy the Redis master. This manifest file specifies a Deployment controller that runs a single replica Redis master Pod:

```

redis-master-deployment.yaml
1 apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
2 kind: Deployment
3 metadata:
4   name: redis-master
5   labels:
6     app: redis
7 spec:
8   selector:
9     matchLabels:
10    app: redis
11    role: master
12    tier: backend
13   replicas: 1
14   template:
15     metadata:
16       labels:
17         app: redis
18         role: master
19         tier: backend
20     spec:
21       containers:
22         - name: master
23           image: k8s.gcr.io/redis:e2e # or just image: redis
24           resources:
25             requests:
26               cpu: 100m
27               memory: 100Mi
28           ports:
29             - containerPort: 6379

```

`kubectl create -f redis-master-deployment.yaml`

- deploy the Redis Master

`kubectl get pods`

- get POD-NAME

```

klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
redis-master-77bb4946d6-w27dz   1/1     Running   0          2m36s
klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$ █

```

`kubectl logs -f POD-NAME` (copy POD-NAME from previous command)

- verify Redis master POD is running

```

[1] 28 Oct 15:03:28.214 # Server started, Redis version 2.8.19
[1] 28 Oct 15:03:28.214 # WARNING you have Transparent Huge Pages (THP) support enabled
# which can cause latency and memory usage issues with Redis. To fix this issue run the command 'echo
# transparent_hugepage/enabled' as root, and add it to your /etc/rc.local in order to retain
# memory settings after THP is disabled.
[1] 28 Oct 15:03:28.215 * The server is now ready to accept connections on port 6379
█

```

hit “**CTRL-C**” to break

Create redis-master service

The guestbook application needs to communicate to Redis master to write it's data. You need to create a Kubernetes Service to proxy the traffic to the Redis master Pod.

A Kubernetes Service is an abstraction which defines a logical set of Pods and a policy by which to access them. It is effectively a named load balancer that proxies traffic to one or more Pods. The Service will tell the Pods to proxy based on Pod labels.

The redis-master-service.yaml manifest file describing a Service resource for the Redis master

```
redis-master-service.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: redis-master
5    labels:
6      app: redis
7      role: master
8      tier: backend
9  spec:
10   ports:
11     - name: redis
12       port: 6379
13       targetPort: 6379
14   selector:
15     app: redis
16     role: master
17     tier: backend
```

This manifest file creates a Service named redis-master with a set of label selectors. These labels match the set of labels that are deployed in the previous step. Therefore, this Service routes the network traffic to the Redis master Pod created in Step 1.

The ports section of the manifest declares a single port mapping. In this case, the Service will route the traffic on port: 6379 to the targetPort: 6379 of the containers that match the specified selector labels. Note that the containerPort used in the Deployment must match the targetPort to route traffic to the Deployment.

```
kubectl apply -f redis-master-service.yaml
- Start the Redis master's Service
```

```
kubectl get service
```

- Verify the service is created

```
klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$ kubectl get service
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
kubernetes     ClusterIP 10.79.240.1 <none>        443/TCP   10m
redis-master   ClusterIP 10.79.243.39 <none>        6379/TCP  73s
klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$
```

## Step 2 – Set up Redis workers

Although the Redis master is a single pod, you can make it highly available and meet traffic demands by adding a few Redis worker replicas.

Take a look at the `redis-slave-deployment.yaml` manifest file describing a Deployment for the Redis worker pods

```
examples/guestbook/redis-slave-deployment.yaml
```

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: redis-slave
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: redis
        role: slave
        tier: backend
    spec:
      containers:
        - name: slave
          image: gcr.io/google_samples/gb-redisslave:v1
          resources:
            requests:
              cpu: 100m
              memory: 100Mi
          env:
            - name: GET_HOSTS_FROM
              value: dns
              # If your cluster config does not include a dns service, then to
              # instead access an environment variable to find the master
              # service's host, comment out the 'value: dns' line above, and
              # uncomment the line below:
              # value: env
      ports:
        - containerPort: 6379
```

Deployment is responsible for achieving the configuration declared in the manifest file. For example, this manifest file defines two replicas for the Redis workers. If there are not any replicas running, Deployment would start the two replicas on your container cluster. If there are more than two replicas are running, it would kill some to meet the specified configuration.

In this case, the Deployment object specifies two replicas. To create the Redis worker Deployment, run:

```
kubectl apply -f redis-slave-deployment.yaml
```

Verify that the two Redis worker replicas are running by querying the list of Pods:

```
kubectl get pods
```

```
klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$ kubectl get pods
NAME                      READY   STATUS    RESTARTS   AGE
redis-master-77bb4946d6-w27dz   1/1     Running   0          9m11s
redis-slave-84548fdb-9ffq9      1/1     Running   0          98s
redis-slave-84548fdb-pq6fq      1/1     Running   0          98s
```

Create redis-slave service

The guestbook application needs to communicate to Redis workers to read data. To make the Redis workers discoverable, you need to set up a Service. A Service provides transparent load balancing to a set of Pods.

The redis-slave-service.yaml defines the Service configuration for the Redis workers

```
redis-slave-service.yaml
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: redis-slave
5    labels:
6      app: redis
7      role: slave
8      tier: backend
9  spec:
10    ports:
11      - port: 6379
12    selector:
13      app: redis
14      role: slave
15      tier: backend
```

This file defines a Service named redis-slave running on port 6379. Note that the selector field of the Service matches the Redis worker Pods created in the previous step.

Create the redis-slave Service by running:

```
kubectl create -f redis-slave-service.yaml
```

```
kubectl get service
```

- verify service is created

```
klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$ kubectl get service
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP  PORT(S)        AGE
kubernetes     ClusterIP  10.79.240.1 <none>       443/TCP       12m
redis-master   ClusterIP  10.79.243.39 <none>       6379/TCP     2m58s
redis-slave    ClusterIP  10.79.240.54 <none>       6379/TCP     8s
klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$
```

### Step 3: Set up the guestbook web frontend

Now that the Redis storage of your guestbook is up, start the guestbook web servers. Like the Redis workers, this is a replicated application managed by a Deployment. This tutorial uses a simple PHP frontend. It is configured to talk to either the Redis worker or master Services, depending on whether the request is a read or a write. It exposes a simple JSON interface, and serves a jQuery-Ajax-based UX. Take a look at the frontend-deployment.yaml manifest file describing the Deployment for the guestbook web server

```
frontend-deployment.yaml
1 apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
2 kind: Deployment
3 metadata:
4   name: frontend
5   labels:
6     app: guestbook
7 spec:
8   selector:
9     matchLabels:
10    app: guestbook
11    tier: frontend
12 replicas: 3
13 template:
14   metadata:
15     labels:
16       app: guestbook
17       tier: frontend
18   spec:
19     containers:
20       - name: php-redis
21         image: gcr.io/google-samples/gb-frontend:v4
22         resources:
23           requests:
24             cpu: 100m
25             memory: 100Mi
26           env:
27             - name: GET_HOSTS_FROM
28               value: dns|
29         ports:
30           - containerPort: 80
```

```
kubectl apply -f frontend-deployment.yaml
```

- create guestbook web frontend Deployment

```
kubectl get pods -l app=guestbook -l tier=frontend
```

- verify the three replicas are running by querying the list of the labels that identify the web frontend (1 of 3 shown running – 2 still being created)

NAME	READY	STATUS	RESTARTS	AGE
frontend-678d98b8f7-5t8gm	0/1	ContainerCreating	0	27s
frontend-678d98b8f7-7zmxf	1/1	Running	0	27s
frontend-678d98b8f7-nx984	0/1	ContainerCreating	0	27s

Expose frontend on an external IP address

The redis-slave and redis-master Services you created in the previous steps are only accessible within the container cluster, because the default type for a Service is ClusterIP.

ClusterIP provides a single IP address for the set of Pods the Service is pointing to. This IP address is accessible only within the cluster.

However, you need the guestbook web frontend Service to be externally visible. That is, you want a client to be able to request the Service from outside the container cluster. To accomplish this, you need to specify type: LoadBalancer in the Service configuration. The frontend-service.yaml manifest file specifying this configuration looks like this

```
frontend-service.yaml
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: frontend
5   labels:
6     app: guestbook
7     tier: frontend
8 spec:
9   type: LoadBalancer
10  ports:
11    - port: 80
12  selector:
13    app: guestbook
14    tier: frontend
```

```
kubectl create -f frontend-service.yaml
```

- create the frontend service

## Step 4: Visit the guestbook website

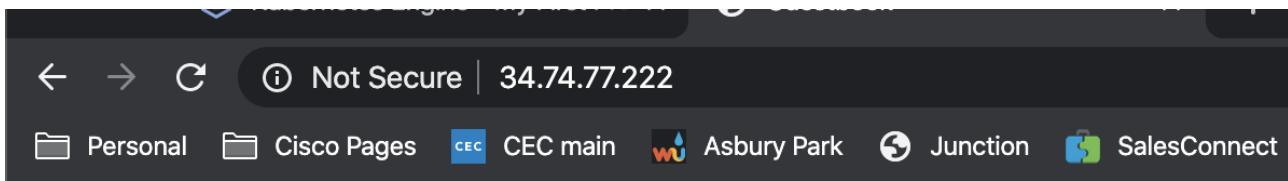
To access the guestbook Service, you need to find the external IP of the Service you just set up by running the command (this may take a minute or so to populate external IP address)

*kubectl get service frontend*

- find the external IP of the Service you just set up (will take a little while)

```
klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$ kubectl get service frontend
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP      PORT(S)      AGE
frontend  LoadBalancer  10.79.242.49  35.185.22.184  80:30343/TCP  47s
klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$
```

Copy the IP address in EXTERNAL-IP column, and load the page in your browser:



## Guestbook

Messages

Submit

type:

*kubectl get pods*

- you should see the 3 frontend PODs we created earlier

```
frontened  Loadbalancer  10.79.242.49  34.74.77.222  80:30343/TCP
kubefy2020@cloudshell:~ (marine-order-267121)$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
frontend-74b4665db5-25w9v  1/1     Running   0          3m28s
frontend-74b4665db5-ptmpn  1/1     Running   0          3m28s
frontend-74b4665db5-qcvkc  1/1     Running   0          3m28s
redis-master-74474c9cd-v66x1 1/1     Running   0          8m22s
redis-slave-74ccb764fc-4xcsv 1/1     Running   0          5m5s
redis-slave-74ccb764fc-cxbp7 1/1     Running   0          5m5s
kubefy2020@cloudshell:~ (marine-order-267121)$
```

Let's scale this up

A ReplicaSet ensures that a specified number of pod replicas are running at any given time. However, a Deployment is a higher-level concept that manages ReplicaSets and provides declarative updates to pods along with a lot of other useful features.

If you look at the ‘frontend-deployment.yaml’ file (3 pages back) you will see we defined 3 replicas in our deployment file.

It is recommended using Deployments instead of directly using ReplicaSets, unless you require custom update orchestration or don’t require updates at all.

For the sake of this lab we will increase the 3 replicas we specified to 5 using the kubectl command.

```
kubectl scale deployment frontend --replicas=5
```

- scale up the number of PODs to 5

```
kubectl get pods
```

You should now see 5 PODs

```
klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$ kubectl get pods
NAME                      READY   STATUS    RESTARTS   AGE
frontend-678d98b8f7-5t8gm  1/1     Running   0          27m
frontend-678d98b8f7-7zmxsf 1/1     Running   0          27m
frontend-678d98b8f7-95vq5  1/1     Running   0          21m
frontend-678d98b8f7-g99gj  1/1     Running   0          21m
frontend-678d98b8f7-nx984  1/1     Running   0          27m
redis-master-77bb4946d6-w27dz 1/1     Running   0          38m
redis-slave-84548fdb-9fffq9 1/1     Running   0          30m
redis-slave-84548fdb-pq6fq  1/1     Running   0          30m
klowenst@cloudshell:~ (gcp-klowenst041320-nprd-68609)$ █
```

## Let's Cleanup

`kubectl delete service frontend`

- delete the service

`gcloud compute forwarding-rules list`

- verify deleted

```
kubeuser@k8s-cisco-lab1:~$ gcloud compute forwarding-rules list
NAME          REGION     IP_ADDRESS   IP_PROTOCOL TARGET
a0a7bce70f21611e7b13c42010a8e01b us-east1  104.196.185.172 TCP      us-east1/targetPools/a0a7bce70f21611e7b13c42010a8e01b
a302e292ff3d711e79cc942010a8e006 us-east1  35.196.99.109  TCP      us-east1/targetPools/a302e292ff3d711e79cc942010a8e006
a5ffa3f17f21611e7b13c42010a8e01b us-east1  35.190.140.202 TCP      us-east1/targetPools/a5ffa3f17f21611e7b13c42010a8e01b
a5b4f32bef17a11e79c1b42010a80009 us-central1 35.226.129.173 TCP      us-central1/targetPools/a5b4f32bef17a11e79c1b42010a80009
a71b31f22f17811e79c1b42010a80009 us-central1 35.225.194.254 TCP      us-central1/targetPools/a71b31f22f17811e79c1b42010a80009
kubeuser@k8s-cisco-lab1:~$ gcloud compute forwarding-rules list
NAME          REGION     IP_ADDRESS   IP_PROTOCOL TARGET
```

`gcloud container clusters delete guestbook`

- delete the container cluster

When CLI commands are finished running you will see guestbook cluster deleted from GCP dash

Now let's delete our original cluster

The screenshot shows the 'Kubernetes clusters' section of the GCP dashboard. A single cluster named 'cluster-1' is listed under the 'Kubernetes clusters' heading. The cluster details are as follows:

Name	Location	Cluster size	Total cores	Total memory	Notifications	Labels
cluster-1	us-east1-b	3	3 vCPUs	11.25 GB		

A red arrow points to the trash can icon in the 'Actions' column for the 'cluster-1' row.

Return to GCP > Compute Engine - delete Centos 7 VM

The screenshot shows the 'VMs' section of the GCP Compute Engine dashboard. A single VM named 'centos-7-1' is listed. The VM details are as follows:

Name	Zone	Recommendation	Internal IP	External IP	Connect
centos-7-1	us-east1-b		10.142.0.2	35.185.78.28	SSH

A red arrow points to the three-dot menu icon in the 'Actions' column for the 'centos-7-1' row. A context menu is open, listing options: Start, Stop, Reset, Delete, New instance group, and View logs.

----- DONE -----

# If you're interested - more k8s hands-on...

Go to Cisco devnet site:

<https://devnetsandbox.cisco.com/RM/Topology?c=0a71d70b-8fa3-4bcb-94dd-6c379dd11c24>

The screenshot shows the DevNet Sandbox - Lab Catalog interface. On the left, there's a sidebar with options like 'Sandbox Labs' (selected), 'Reservations' (0 New, 0 Total), and filters for 'Sandbox Lab Status' (Available, Unavailable, View only). The main area is titled 'CLOUD (3)' and shows three lab cards:

- Cisco AppDynamics** (Version 1.0): Cisco AppDynamics Enterprise/Cloud Applications Monitoring & Analytics. Includes a 'RESERVE' button.
- Cisco Container Platform** (Version 5.1): Container management for a multicloud world. Includes a 'RESERVE' button.
- Istio** (Version 1.2): Istio 1.2 with your own Kubernetes 1.15 cluster. Includes a 'RESERVE' button.

A red oval highlights the three lab cards. A sidebar on the right lists categories like Cloud, ALL CATEGORIES, and others.